



# Memory virtualization overhead mitigation using contiguous memory virtual machines

Peterson Yuhala

Supervised by : Prof. Alain Tchana

University of Neuchâtel - June 2018

## Introduction

## Main Issue

Memory address  
translation

## Paging

## Address Translation

## Shadow Paging

## Nested Paging

## Our Solution


## Proof of Concept

## Implementation

## Contributions

## Conclusion

- ▶ Address translation represents a significant part of virtualization overhead [ISCA'16,ISCA'17].
- ▶ Our main goal is to nullify that overhead in virtualized systems.



3/32

- Memory virtualization depends on VM type. Three main mechanisms exist.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

## Introduction

### Main Issue

### Memory address translation

#### Paging

#### Address Translation

### Shadow Paging

### Nested Paging

### Our Solution

### Proof of Concept

### Implementation

### Contributions

### Conclusion

- ▶ Memory virtualization depends on VM type. Three main mechanisms exist.
- ▶ In paravirtualised systems, VMs use a technique called **direct paging** [Xen'15].
- ▶ Here the guest uses a **Physical to Machine** mapping maintained within the hypervisor to map virtual addresses directly to physical addresses in its page table [Xen'15].

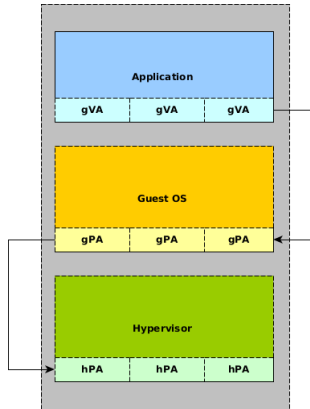
- ▶ Memory virtualization depends on VM type. Three main mechanisms exist.
- ▶ In paravirtualised systems, VMs use a technique called **direct paging** [Xen'15].
- ▶ Here the guest uses a **Physical to Machine** mapping maintained within the hypervisor to map virtual addresses directly to physical addresses in its page table [Xen'15].
- ▶ For hardware assisted virtualised systems, we have two main techniques used: **shadow paging** and **nested paging**.

- ▶ Memory virtualization depends on VM type. Three main mechanisms exist.
- ▶ In paravirtualised systems, VMs use a technique called **direct paging** [Xen'15].
- ▶ Here the guest uses a **Physical to Machine** mapping maintained within the hypervisor to map virtual addresses directly to physical addresses in its page table [Xen'15].
- ▶ For hardware assisted virtualised systems, we have two main techniques used: **shadow paging** and **nested paging**.
- ▶ Our main focus here is hardware assisted virtualised systems as HVMs are becoming the norm [Amazon AWS].

- The main issue here is **Memory Address Translation** in virtualized systems. ie Obtaining the corresponding **machine address** in host address space given a **virtual address** in guest address space.



- Main memory addresses are called **host physical addresses** (hPA) and process addresses are called **guest virtual addresses** (gVA). Addresses in the guest address space are referred to as **guest physical addresses** (gPA).



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution


Proof of Concept

Implementation

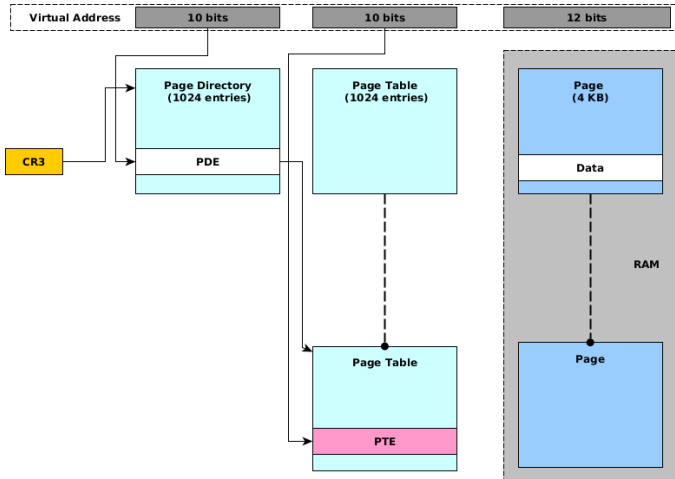
Contributions

Conclusion

- ▶ Native systems (No Virtualization) require one-level address translations. ie **virtual address** → **physical address**.



- The following figure illustrates the page walking mechanism in a two-level page table.



7/32

32

Introduction

Main Issue

Memory address translation

Paging

Address Translation

8/32

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- Virtualized systems on the other hand require two-level address translations. ie  $gVA \rightarrow gPA$  AND  $gPA \rightarrow hPA$ .



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

9/32

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- Two approaches are used : **Shadow Paging** and **Nested Paging**.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

10/32

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- ▶ In shadow paging the hypervisor maintains a **shadow page table** for **gVA**  $\rightarrow$  **hPA** mappings and the **CR3** register points to the shadow page table.
- ▶ The hypervisor captures any attempts by the VM to update its page tables and uses **gVA**  $\rightarrow$  **gPA** and **gPA**  $\rightarrow$  **hPA** to construct the shadow page table.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

10/32

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- ▶ In shadow paging the hypervisor maintains a **shadow page table** for **gVA** → **hPA** mappings and the **CR3** register points to the shadow page table.
- ▶ The hypervisor captures any attempts by the VM to update its page tables and uses **gVA** → **gPA** and **gPA** → **hPA** to construct the shadow page table.
- ▶ On a TLB miss, the hardware performs a 1D page walk on the shadow page table to obtain the machine address corresponding to a guest virtual address.

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

**Shadow Paging**

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

11/32

- Only 4 memory references (same as base native) required for a complete page walk on a TLB miss [ISCA'16].



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

**Shadow Paging**

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

11/32

- ▶ Only 4 memory references (same as base native) required for a complete page walk on a TLB miss [ISCA'16].
- ▶ No need for any extra hardware support for page walks.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

**Shadow Paging**

12/32

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- ▶ Every page table update requires a costly trap in the hypervisor to update shadow page table entries [ISCA'16].

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

12/32

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- ▶ Every page table update requires a costly trap in the hypervisor to update shadow page table entries [ISCA'16].
- ▶ Costly traps in the hypervisor on context switches [ISCA'16].

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

**Nested Paging**

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

13/32

- ▶ Nested paging is a widely used hardware technique to virtualize memory and here the MMU a **two dimensional page table** consisting of the **guest page table** (gPT) and **host page table** (hPT).

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

**Nested Paging**

Our Solution

Proof of Concept

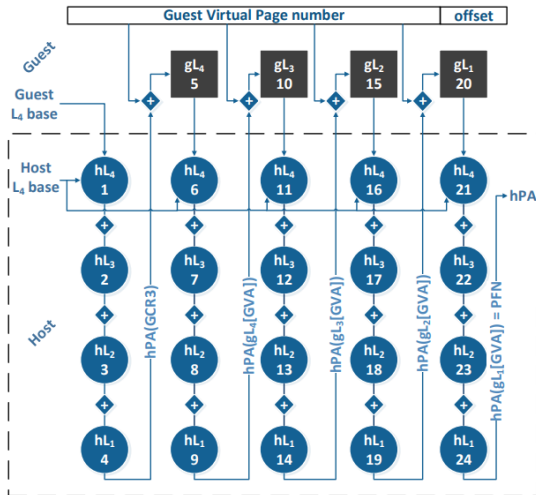
Implementation

Contributions

Conclusion

13/32

- ▶ Nested paging is a widely used hardware technique to virtualize memory and here the MMU a **two dimensional page table** consisting of the **guest page table** (gPT) and **host page table** (hPT).
- ▶ On a TLB miss the processor performs a **2D page walk** for any **gVA** → **hPA** translation.



Figure

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

**Nested Paging**

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

15/32

- Allows fast direct updates to both page tables (gPT and hPT) without any traps in the hypervisor (ie no VMM intervention).

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

**Nested Paging**

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

16/32

- On a TLB miss the MMU performs a 2D page walk that requires up to 24 memory references. Figure 1 shows how up to 24 memory references can be obtained for a complete 2D page walk.



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

17/32

	Base Native	Nested Paging	Shadow Paging
<b>TLB Hit</b>	fast	fast	fast
<b>Max. memory access on TLB miss</b>	4	24	4
<b>Page table updates</b>	fast direct	fast direct	slow mediated by VMM
<b>Hardware support</b>	1D page walk	2D page walk	1D page walk

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

**Our Solution**

Proof of Concept

Implementation

Contributions

Conclusion

18/32

- From the table, Nested and Shadow paging both present significant overhead due to memory accesses and page table updates respectively.

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

**Our Solution**

Proof of Concept

Implementation



Contributions

Conclusion

18/32

- ▶ From the table, Nested and Shadow paging both present significant overhead due to memory accesses and page table updates respectively.
- ▶ We propose a new mechanism of address translation which maintains the advantages of both Nested and Shadow paging, and at the same time nullifies the overhead of both mechanisms.

- Our basic idea is to give **contiguous machine memory** to VMs and introduce a hardware feature to store the **base address** of the **contiguous chunk**.



19/32

32

Introduction

Main Issue

Memory address translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept



Implementation

Contributions

Conclusion

19/32

- Our basic idea is to give **contiguous machine memory** to VMs and introduce a hardware feature to store the **base address** of the **contiguous chunk**.



- Using our solution, the **gPT** and the **base address** are sufficient for a complete **gVA** → **hPA** translation.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

**Our Solution**

Proof of Concept

Implementation

Contributions

Conclusion

20/32

	Base Native	Nested Paging	Shadow Paging	Our Solution
<b>TLB Hit</b>	fast	fast	fast	fast
<b>Max. memory access on TLB miss</b>	4	24	4	4
<b>Page table updates</b>	fast direct	fast direct	slow mediated by VMM	fast direct
<b>Hardware support</b>	1D page walk	2D page walk	1D page walk	1D page walk

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

Implementation

Contributions

Conclusion

21/32

- ▶ We made a **simulator** to show that it is possible to give contiguous memory to a high percentage of VMs in a datacenter.
- ▶ We simulated VM traces from Microsoft Azure [SOSP'17] and Bitbrains.
- ▶ The VM traces contain information about VM creation and destruction timestamps, CPU utilization and VM memory.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

22/32

Implementation

Contributions

Conclusion

- ▶ The simulator simulates 2 VM placement algorithms :
  - ▶ **Traditional placement algorithm** which takes into account resource availability (spread and stack).
  - ▶ **Contiguity-Aware placement algorithm** which prioritizes servers which can provide contiguous memory for VMs.



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

Implementation

Contributions

Conclusion

22/32

- ▶ The simulator simulates 2 VM placement algorithms :
  - ▶ **Traditional placement algorithm** which takes into account resource availability (spread and stack).
  - ▶ **Contiguity-Aware placement algorithm** which prioritizes servers which can provide contiguous memory for VMs.
- ▶ It takes as input VM traces and then calculates the percentage of VMs with contiguous memory at the end of the simulation.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

Implementation

Contributions

Conclusion

23/32

- ▶ We use different generations of MS Azure servers in our simulator. ie. generation 3 to 6 and HPC servers so as to reflect the reality in real datacenters.
- ▶ A mix of servers with the following composition was used in our simulator :

Gen.	%
Gen3	20
HPC	20
Gen4	20
Gen5	20
Gen6	10
Godzilla	10

Introduction

Main Issue

Memory address translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

Implementation

Contributions

Conclusion

24/32

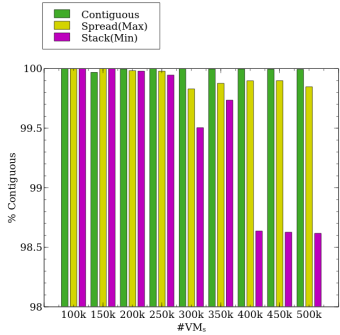


Figure : % contiguous-memory VMs using different placement algorithms

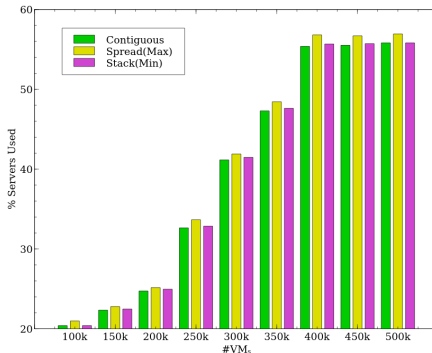


Figure : Server consolidation rate for both contiguous-aware and traditional placement algorithms.

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

**Proof of Concept**

Implementation

Contributions

Conclusion

26/32

- ▶ We have a high % of VMs with contiguous memory because VMs sizes are relatively small compared to server sizes [SOSP'17].
- ▶ The relatively small and discrete size of VMs reduces the impact of memory fragmentation and prevents resource sprawl in the datacenter.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

**Implementation**

Contributions

Conclusion

27/32

- We implement our solution in Xen hypervisor. Here we modify the default Xen memory allocator so it allocates contiguous chunks of machine memory to VMs.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

27/32

- ▶ We implement our solution in Xen hypervisor. Here we modify the default Xen memory allocator so it allocates contiguous chunks of machine memory to VMs.
- ▶ To provide contiguous memory to a VM in Xen, we calculate the **order** corresponding to the VM's memory demand. ie **VM Total Memory =  $2^{\text{order}} * 4 \text{ KB}$**  .
- ▶ This order is then passed to the **Xen heap allocator** which provides a contiguous chunk of memory containing  $2^{\text{order}}$  contiguous 4kb pages.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

**Implementation**

Contributions

Conclusion

- We succeed already in giving contiguous memory to PV VMs. We would then eliminate unnecessary hypercalls knowing that VMs have contiguous memory.

28/32



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

**Implementation**

Contributions

Conclusion

28/32

- ▶ We succeed already in giving contiguous memory to PV VMs. We would then eliminate unnecessary hypercalls knowing that VMs have contiguous memory.
- ▶ We succeed in mapping 99% of HVM memory to contiguous machine memory. We would modify the Xen heap allocator to give 100% contiguity. The memory translation algorithm would then be modified as per our basic idea.

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

**Implementation**

Contributions

Conclusion

28/32

- ▶ We succeed already in giving contiguous memory to PV VMs. We would then eliminate unnecessary hypercalls knowing that VMs have contiguous memory.
- ▶ We succeed in mapping 99% of HVM memory to contiguous machine memory. We would modify the Xen heap allocator to give 100% contiguity. The memory translation algorithm would then be modified as per our basic idea.
- ▶ The last step would be the implementation of the whole solution in a simulator and processing of final results.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

- **Memory isolation** : Each VM must be confined within its memory region. How do we ensure isolation knowing that the hypervisor has minimal interference.

29/32

32

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

29/32

- ▶ **Memory isolation** : Each VM must be confined within its memory region. How do we ensure isolation knowing that the hypervisor has minimal interference.
- ▶ How to **enforce contiguity**.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

29/32

- ▶ **Memory isolation** : Each VM must be confined within its memory region. How do we ensure isolation knowing that the hypervisor has minimal interference.
- ▶ How to **enforce contiguity**.
- ▶ **Cohabitation** of our solution with Nested Paging or Shadow Paging. ie Contiguous memory VMs use our solution and other VMs use Nested Paging/Shadow Paging.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

**Contributions**

Conclusion

30/32

### Memory Isolation :

- Introduction of a base address (**gBA**) for each VM (included within the VM Control Structure (VMCS)).

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

**Contributions**

Conclusion

30/32

### Memory Isolation :

- ▶ Introduction of a base address (**gBA**) for each VM (included within the VM Control Structure (VMCS)).
- ▶ Once scheduled, the VM's **gBA** and VM's memory size (**gMS**) are stored within dedicated processor registers.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

30/32

### Memory Isolation :

- ▶ Introduction of a base address (**gBA**) for each VM (included within the VM Control Structure (VMCS)).
- ▶ Once scheduled, the VM's **gBA** and VM's memory size (**gMS**) are stored within dedicated processor registers.
- ▶ To ensure isolation, the TLB is enhanced with a **security module** such that during insertion into the TLB, the security module checks if  $hPA \leq gBA + gMS$



Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

**Contributions**

Conclusion

31/32

## Memory Contiguity :

- At boot time, we allocate contiguous memory to VMs if possible.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

31/32

## Memory Contiguity :

- ▶ At boot time, we allocate contiguous memory to VMs if possible.
- ▶ For VMs with discontinuous memory, the hypervisor reorganizes the VM's memory to be contiguous. This can be done either by migrating the VM's pages during idle periods or by migrating the VM to another host.

- Memory virtualization overhead can be very high when using traditional techniques (Nested and Shadow Paging).

Introduction  
Main Issue  
Memory address translation  
    Paging  
    Address Translation  
Shadow Paging  
Nested Paging  
Our Solution  
Proof of Concept  
Implementation  
Contributions  
Conclusion

32/32

Introduction

Main Issue

Memory address  
translation

Paging  
Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

32/32

- ▶ Memory virtualization overhead can be very high when using traditional techniques (Nested and Shadow Paging).
- ▶ We propose to eliminate the overhead of these traditional techniques using contiguous-memory VMs.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

32/32

- ▶ Memory virtualization overhead can be very high when using traditional techniques (Nested and Shadow Paging).
- ▶ We propose to eliminate the overhead of these traditional techniques using contiguous-memory VMs.
- ▶ Through simulation we showed that it is possible to give contiguous memory to a high % of VMs in datacenters.

Introduction

Main Issue

Memory address  
translation

Paging

Address Translation

Shadow Paging

Nested Paging

Our Solution

Proof of Concept

Implementation

Contributions

Conclusion

32/32

- ▶ Memory virtualization overhead can be very high when using traditional techniques (Nested and Shadow Paging).
- ▶ We propose to eliminate the overhead of these traditional techniques using contiguous-memory VMs.
- ▶ Through simulation we showed that it is possible to give contiguous memory to a high % of VMs in datacenters.
- ▶ Our solution is being prototyped in the Xen hypervisor; we would then evaluate its performance to show that it is indeed better than Nested and Shadow paging.

Thank you for listening !

