

# Remove hardware-assisted page walk overhead for virtualized systems using memory contiguity property of VMs

---

Alain Tchana & Boris Teabe

Institut National Polytechnique de Toulouse - IRIT

2018

# Plan

---

1 Memory address translation

2 Contributions:

# Sous-section 1

---

1 Memory address translation

2 Contributions:

# Paging

---

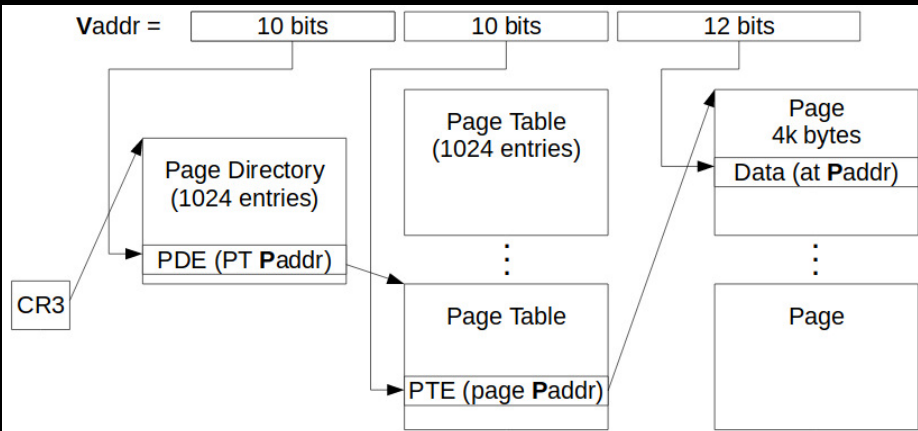
## *Principle*

---

- ▶ Each process has a virtual address space
  - ▶ 0 to  $2^{32}$  (in 32-bit machines)
  - ▶ 0 to  $2^{64}$  (in 64-bit machines)
- ▶ Main memory addresses are called host physical address (hPA)
- ▶ Process addresses are called guest virtual address (gVA)
- ▶ Each process address space is organized in memory pages
  - ▶ 4KB, 2MB (large page)

# Native systems

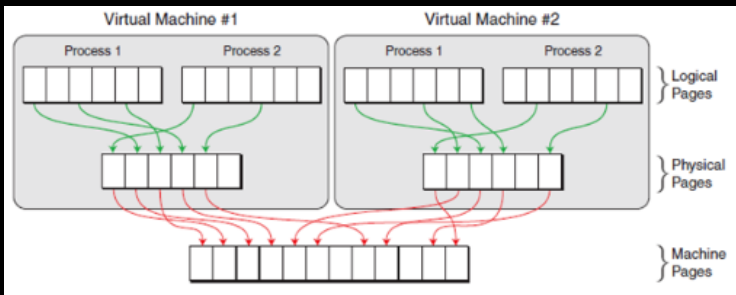
$gVA \rightarrow hPH$



# Virtualized systems

## *Address Translation for VMs*


- ▶ Guest virtual address (gVA)  $\rightarrow$  guest physical address (gPA)  $\leftarrow$  host physical address (hPA)
- ▶ Two approaches: Shadow paging and Hardware assisted virtualization



# Virtualized systems

## Shadow Paging

- ▶ Hypervisor maintains a separate mapping from gVA  $\rightarrow$  hPA
  - ▶ Hypervisor intercepts every attempt by guest to update or install a page table (performance)
  - ▶ Per process mapping (space)
- ▶ This approach incurs lots of page faults and VM exits



# Virtualized systems

---

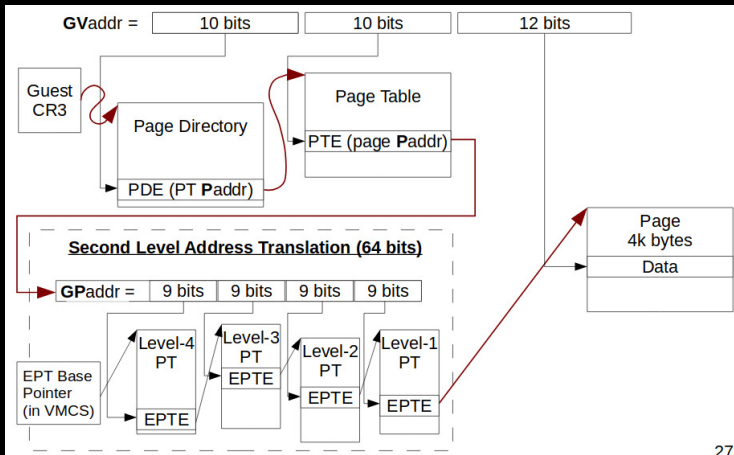
## *Hardware assisted virtualization*

- ▶ This approach eliminates the need for shadow paging
- ▶ And provides architectural support for a new layer of address translation (also called Nested Level)
- ▶ A new hierarchy of paging: gPA  $\rightarrow$  gPA
- ▶ Two dimensional Page Table
  - ▶ Nested Page Table (NPT) by AMD or Extended Page Table (EPT) by Intel



# Virtualized systems

## Hardware assisted virtualization: $gPA \rightarrow hPA$



27

# Virtualized systems

## *Hardware assisted virtualization: Overhead*

- ▶ Each TLB miss leads to 24 memory access
  - ▶ Guest page table walk
  - ▶ Nested or Extended page table walk
- ▶ TLB miss intensive applications are penalized, in comparison to software based MMU (Shadow paging)


	Base Native	Nested Paging	Shadow Paging
TLB hit	fast (VA $\Rightarrow$ PA)	fast (gVA $\Rightarrow$ hPA)	fast (gVA $\Rightarrow$ hPA)
Max. memory access on TLB miss	4	24	4
Page table updates	fast direct	fast direct	slow mediated by VMM
Hardware support	1D page walk	2D+1D page walk	1D page walk

# Virtualized systems

## *Hardware assisted virtualization: Overhead*

from [https://www.vmware.com/pdf/Perf\\_ESX\\_Intel-EPT-eval.pdf](https://www.vmware.com/pdf/Perf_ESX_Intel-EPT-eval.pdf)

- Higher is Better



# Related work

---

## *Existing solutions*

---

- ▶ To increase TLB size, using the CPU cache (ISCA '17)
- ▶ Direct-segment (ISCA '12, ISCA '17)
- ▶ Large pages (ASPLOS '08)
- ▶ Caching nested page walk (MICRO '14)
- ▶ Switch between Nested or Extended page table  $\Leftrightarrow$  Shadow page table (ISCA '14)

# Sous-section 2

---

① Memory address translation

② Contributions:

# Basic idea

---

## *Single page table walk*


---

- ▶ Remove EPT
- ▶ The guest's page table is sufficient when the host memory of a VM is contiguous
- ▶  $gVA \rightarrow hPH$  is computed as follows
  - ▶  $gVA \rightarrow gPA$
  - ▶  $gPA + gBaseAddr = hPA$
  - ▶  $gBaseAddr$  is the host address start for the VM in the host memory
- ▶ This idea is inspired by "*Do-It-Yourself Virtual Memory Translation*", ISCA'17

# Basic idea

## *Single page gable walk*

- ▶ For a contiguous memory VM:
  - ▶ The VM OS is in charge of handling his page tables without hypervisor interfering
  - ▶ Single page walk after a TLB miss



# A favorable context: Contiguous memory VM

## *Existence of VMs with contiguous Memory*

- ▶ Simulate the hypervisor VM memory allocator at boot (Xen hypervisor)
  - ▶ allocates contiguous memory region
- ▶ Simulate two memory placement algorithms for VMs
  - ▶ naive placement which only takes into account resource availability
  - ▶ contiguity aware placement which prioritises machines with contiguous memory regions
- ▶ Run the simulation on Microsoft Azure traces
  - ▶ take into account VM arrival rate
  - ▶ count the number of VMs having contiguous memory



# Challenges

---

## *Memory isolation*

---

- ▶ ( $C_1$ ) Each VM must be confined within its memory region
  - ▶ How to ensure isolation, knowing the hypervisor does not interfere?
- ▶ ( $C_2$ ) How to enforce contiguity?
- ▶ ( $C_3$ ) Cohabitation of our solution with EPT
  - ▶ Contiguous memory VMs use our solution
  - ▶ Other VMs use EPT


# Contributions

## *Memory isolation*

- ▶ Introducing a base address (noted **gBA**) for each VM ( included within the VMCS)
- ▶ Once scheduled, the VM's **gBA** is loaded within a dedicated processor register (guest base address register).
- ▶ Also, the VM's memory size (noted **gMS**) is stored in a dedicated processor register .
- ▶ To ensure isolation (no memory violation), the TLB is enhanced with a security module
  - ▶ during insertion in TLB, the security module checks if  $\text{gBA} \leq \text{hPA} \leq \text{gBA} + \text{gMS}$

# Contributions

## Memory isolation



# Contributions

---

## *Memory contiguity enforcement: in the datacenter*

- ▶ We improve the VM placement system within the datacenter in order to prioritize machines offering contiguity to the VM
- ▶ In addition to classical monitoring informations (CPU, memory), each compute node sends the size of the contiguous memory regions
- ▶ Placement system rely on these information to place

# Contributions

---

## *Memory contiguity enforcement: on a server*

---

- ▶ We change a VM from discontinuous memory to contiguous memory at runtime
  - ▶ The VM boot with discontinuous memory
  - ▶ Once possible, the hypervisor reorganizes the VM's memory to be contiguous
  - ▶ This can be done locally (migrating the VM's pages during idle period), or migrate the VM to another host