



Out of Hypervisor (OoH): Efficient Dirty Page Tracking In Userspace Using Hardware Virtualization Features

Stella Bitchebe - (bitchebe@i3s.unice.fr)

Advisor: Pr Alain Tchan - (alain.tchan@ens-lyon.fr)

ENS Lyon

Rennes - April, 29

1 Dirty Page Tracking

2 OoH for PML

Dirty Page Tracking in Userspace

Purpose

- ▶ WSS (working set size) estimation
- ▶ Live migration
- ▶ Checkpointing
- ▶ Garbage collection

Current approach

- ▶ Page write protection
- ▶ Two main solutions
 - ▶ Linux /proc interface
 - ▶ Linux userfaultfd (ufd) interface

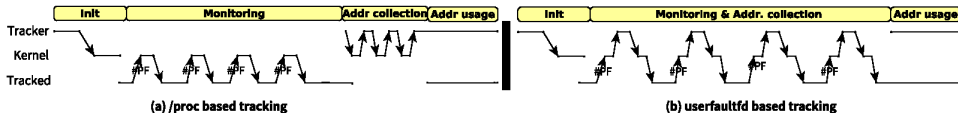
Dirty Page Tracking in Userspace

Nomenclature

- ▶ Tracker: the monitoring thread (e.g., CRIU, Boehm GC)
- ▶ Tracked: the thread whose memory is monitored (any application)

Overall Functioning

- ▶ Tracker's activity can be organized in four phases:
 - ▶ the initialization of the tracking method,
 - ▶ the monitoring
 - ▶ the collection of dirty page addresses
 - ▶ the exploitation of the latter (e.g., for checkpointing)



Page write protection

Overhead

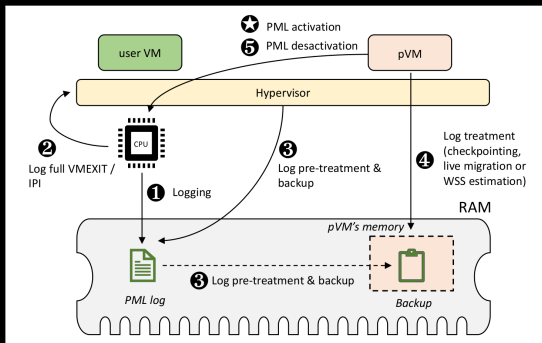
- ▶ `ufd`
 - ▶ We measured up to $15.6\times$ and $14.5\times$ slowdown for 1GB on Tracked and Tracker respectively
 - ▶ Due to page fault handling and context switches
- ▶ `/proc`
 - ▶ We measured about 2.234ms to parse the PT and flush the TLB (in the kernel), and up to $4.3\times$ and $2.5\times$ slowdown for 1GB on Tracked and Tracker respectively
 - ▶ We measured about 594.187ms to parse the PT in userspace (`/proc/PID/pagemap`) for 1GB

1 Dirty Page Tracking

2 OoH for PML

Intel PML

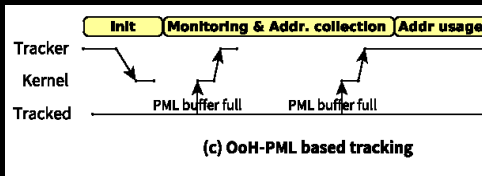
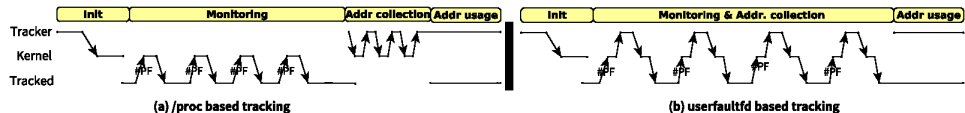
Functioning



Intel PML in the OoH Context

- To accelerate CRIU checkpointing and Boehm garbage collection

Dirty page tracking in userspace (with PML)



OoH for PML

Challenges

- ▶ (C_1) PML can only be managed by the hypervisor
- ▶ (C_2) PML works at coarse-grained, that is it concerns the entire VM
- ▶ (C_3) PML only logs GPAs

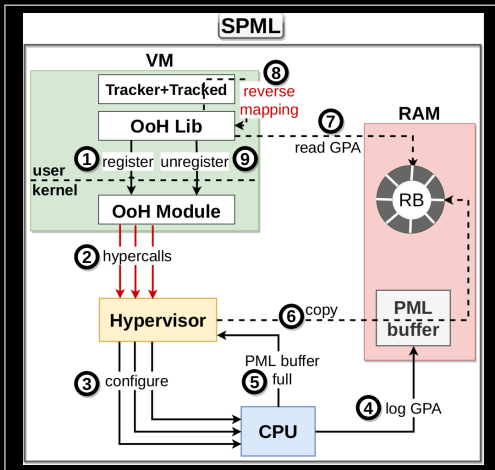
OoH for PML

Two Solutions

- ▶ Extended PML (EPML): small hardware changes
- ▶ Shadow PML (SPML): no hardware modification
 - ▶ Its huge overhead justifies EPML

OoH for PML

SPML Design



OoH for PML

SPML limitations

- Costly reverse mapping (takes about 15.739 s for 1GB working set)

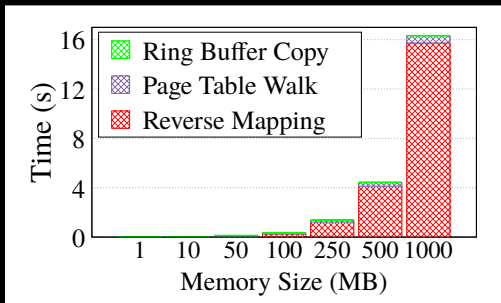


Figure: Reverse mapping appears to be the bottleneck of SPML.

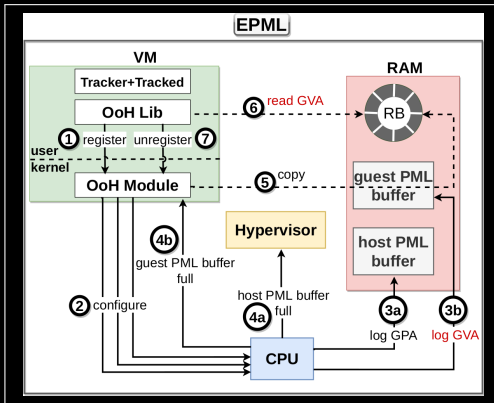
OoH for PML

SPML limitations

- ▶ Costly reverse mapping (takes about 15.739 s for 1GB working set)
- ▶ Costly hypercalls ($4.49\mu s$ to perform an empty hypercall)

OoH for PML

EPML Design



OoH for PML

EPML

- ▶ We leverage VMCS shadowing to allow the guest to perform `vmread` and `vmwrite` instructions
- ▶ At load time, OoH Module calls the hypervisor to enable and configure VMCS shadowing

OoH for PML

Hardware changes

- ▶ We introduce in VMCS's VM-Execution Control area a new field (called Guest PML Address)
- ▶ We extend the PT walk to make the processor log GVA to the guest-level PML buffer and the GPA to the hypervisor-level PML buffer.
- ▶ We modify the hardware so that when the guest-level PML buffer is full, the processor raises a virtual self-IPI (Inter-Processor Interrupt)

OoH for PML

Implementation

- ▶ We implemented EPML's hardware changes in BOCHS
- ▶ We used Xen as the hypervisor and Linux as the guest OS
- ▶ We integrated OoH Lib with CRIU and Boehm GC

OoH for PML

Benchmarks

- ▶ Macro-benchmarks: tkrzw **tkrzw** applications (key value store) and Phoenix applications (MapReduce)
- ▶ We considered three working set sizes (Small, Medium and Large)

OoH for PML

Evaluations

- ▶ What is the potential overhead or improvement of SPML and EPML compared to `/proc` and `ufd`?
- ▶ What is the scalability of SPML and EPML?
- ▶ How to accurately evaluate EPML?
- ▶ Approach
 - ▶ build a formula
 - ▶ show the accuracy of that formula on other techniques, which are measurable
 - ▶ by construction, the formula is accurate for EPML

Evaluations

- ▶ The execution time of Tracker when it implements technique x

$$E(C_{t\text{ker}}) = E(C_x) + E(C_p) + I(C_x, C_p) \quad (1)$$

where $E(C)$ is the execution time of code C and $I(C_1, C_2)$ is the impact of C_1 on C_2 .

Evaluations

Impact on Tracker

$$\begin{aligned} E(C_{/proc}) &= E(C_{echo\ 4\ >\ /proc/PID/clear_refs}) \\ &\quad + E(C_{page\ table\ walk\ in\ userspace}) \\ E(C_{UFD}) &= E(C_{ioctl\ write_protect}) \\ &\quad + E(C_{ioctl\ register}) \\ &\quad + E(C_{ioctl\ write_unprotect}) \\ E(C_{SPML}) &= E(C_{ring\ buffer\ copy}) \\ &\quad + E(C_{reverse\ mapping}) \\ &\quad + E(C_{enable/disable\ PML}) \\ E(C_{EPML}) &= E(C_{ring\ buffer\ copy}) \\ &\quad + E(C_{enable/disable\ PML}) \end{aligned} \tag{2}$$

Evaluations

Evaluations

- The execution time of Tracked when it is monitored by a tracker using the technique x :

$$E(C_{tked_tker}) = E(C_{tked}) + E(C_{tker}) + I(C_x, C_{tked}) \quad (3)$$

where $I(C_x, C_{tked})$ consists of page faults, vmexits, etc.. Thus, the overhead of x on Tracked is $E(C_{tker}) + I(C_x, C_{tked})$.

Evaluations

Impact on Tracked

$$\begin{aligned} I(C_{/proc}, C_{tked}) &= E(C_{PFH \text{ in kernelspace}}) \\ &\quad + E(C_{context \text{ switch}}) \\ I(C_{UFD}, C_{tked}) &= E(C_{PFH \text{ in userspace}}) \\ &\quad + E(C_{context \text{ switch}}) \\ I(C_{SPML}, C_{tked}) &= E(C_{vmexits}) \\ &\quad + N \times E(C_{vmread/vmwrite}) \\ I(C_{EPML}, C_{tked}) &= N \times E(C_{vmread/vmwrite}) \end{aligned} \tag{4}$$

For $I(C_{EPML}, C_{tked})$, we use SPML's N (the number of context switches) as it is the same (validated by running SPML and EPML under BOCHS).

Formula validation

Metric	Time (ms)
$E(C_{tker})$ measured	5503.79
$E(C_{tked_tker})$ measured	135255.35
$E(C_p)$	251.35
$E(C_{copy_rb})$	0.49
$E(C_{disable\ pml})$	2.06
$E(C_{rev.\ mapping})$	5419
$E(C_{tker})$ estimated	5672.9
$E(C_{vmexits})$	18000
N	39
$E(C_{vmread,vmwrite})$	1.73×10^{-3}
$E(C_{tked_tker})$ estimated	136919.85

(a) SPML

Metric	Time (ms)
$E(C_{tker})$ measured	1097.99
$E(C_{tked_tker})$ measured	115283.98
$E(C_p)$	251.35
$E(C_{clear_refs})$	1.409
$E(C_{PTwalk})$	0.89
$E(C_{tker})$ estimated	1116.09
$E(C_{PFHuser})$	0.27
$E(C_{tked_tker})$ estimated	114418.58

(b) /proc

Evaluations

Formula validation

An accuracy of about 96.34% and 99% respectively for SPML and /proc formulas

Evaluations

Impact on Boehm

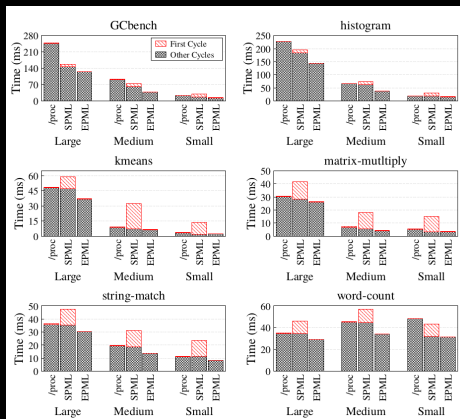


Figure: We highlight the first GC cycle during which Boehm performs reverse mapping with SPML.

Evaluations

Impact on applications

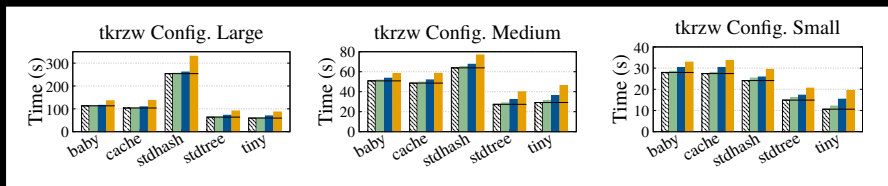


Figure: Impact of Boehm GC.

Evaluations

Impact on CRIU

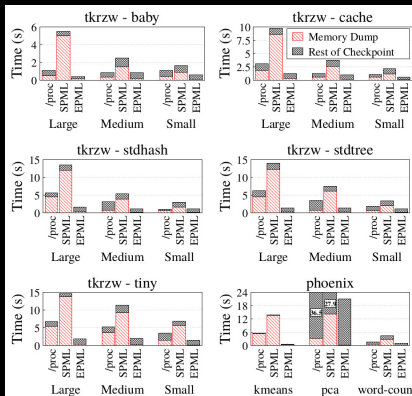
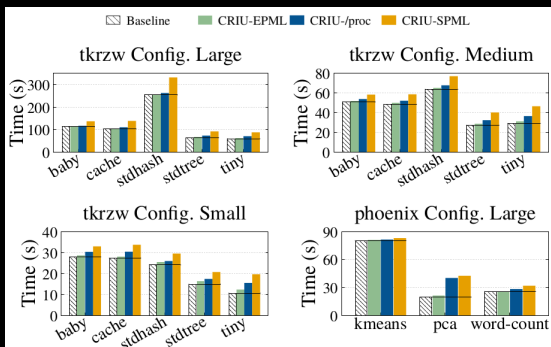


Figure: We highlight MD phase during which CRIU performs reverse mapping with SPML.

Evaluations

Impact on the checkpointed app



Evaluations

Summary

- ▶ Impact on the tracker
 - ▶ SPML induces up to $5\times$ slowdown on CRIU and $3\times$ slowdown on Boehm GC.
 - ▶ EPML brings up to $4\times$ speedup compared to /proc and $13\times$ speedup compared to SPML for CRIU; up to $2\times$ speedup compared to /proc and up to $6\times$ speedup compared to SPML.

Evaluations

Summary

- ▶ Impact on the application performance
 - ▶ `/proc`, which is the default solution implemented in both CRIU and Boehm, incurs an overhead of up to 102% with CRIU on the Phoenix `pca` application, and up to 232% with Boehm on the Phoenix `string-match` application.
 - ▶ The overhead of SPML is up to 114% with CRIU and 273% with Boehm on the same applications.
 - ▶ EPML leads to the lowest overhead, which is about 7% with CRIU and 24% with Boehm.

Summary

Dirty Page Tracking and OoH for Intel PML

- ▶ For wss estimation, live migration, checkpointing, GC, ...
- ▶ Induce high overhead on applications
- ▶ For improving process/container checkpointing, concurrent GCs
- ▶ Excellent results in both reducing overhead on tracked applications and improving the tracking application

Take Away

- ▶ Existing software-based tools can be improved using hardware virtualization features
- ▶ When thinking hypervisor-oriented hardware virtualization features, we must think how they can also be used by processes from inside VMs, this may need few efforts